

# Outline



**Lattice  
Agreement**

Ordering  
issues



Generalized Lattice  
Agreement

Nullifying  
semantics

Sagar Chordia  
With {Kapil, Rama, Kaushik}  
PLATO - MSRI


Commutative  
Replicated Data-  
types

Applications

# Lattice agreement

- $N$  processes.
- Each process starts with an input value  $u_i \in (L, \leq, V)$  a finite join semi lattice.
  - $\leq$  is partial order
  - For any two elements  $a, b \in L$ .  $a \vee b$  exists.
- Every non-faulty process outputs a value  $v_i$ 
  - $v_i$  is join of input values including its own.
  - Any two output values  $v_i$  and  $v_j$  are comparable i.e. either  $v_i \leq v_j$  or  $v_j \leq v_i$
  - Every correct process eventually outputs a value.

# Lattice agreement



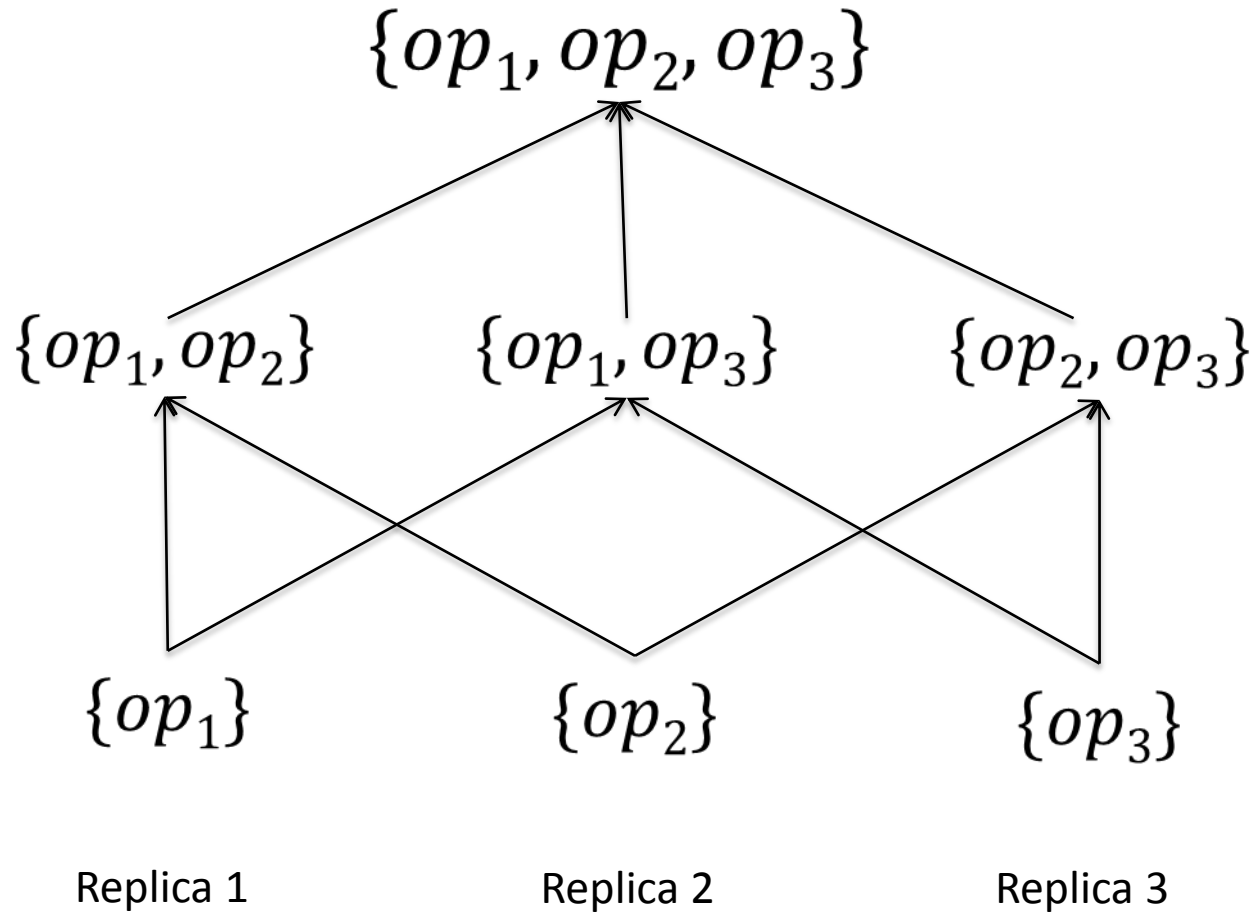
# Generalized lattice agreement

- Generalization of lattice agreement
  - Process receives sequence of values  $u_i^j$
  - Values belong to an infinite lattice.
- Processes output sequence of values  $v_i^j$ 
  - Any two values are comparable.
  - Every value received by correct process is eventually included in an output value.
- Wait-free algorithm.  $O(n)$  message delays.

# State machine replication

- GLA can be used to implement a specific class of state machines
  - Two kinds of operations (a) updates, and (b) reads
  - All concurrent updates should *commute*
- Translation to GLA
  - Updates form a power-set lattice  $(L, \leq, \vee)$
  - Replicas propose operations
  - Queries are serviced using states corresponding to output values.

# SM with GLA




# Guarantees

- Fault-tolerant model.
- Wait-free algorithm.  $O(n)$  message delays.
- Strong consistencies
  - Sequential consistency
    - Serializability and program order of client operations is maintained.
  - Linearizability
    - Operations appear to execute instantaneously.

# Data-types with GLA as SM


- Reads and updates are only operations.
- Updates are commutative
- Set with `add()` as only update operation.





# Queries?

- Even if updates commute, queries may be inconsistent



- As GLA learns chain of values

$\{\}, \{2\}, \{1, 2\}$  ✓


$\{\}, \{1\}, \{1, 2\}$  ✓

# Commutative and Convergent Replicated Data-Types



- Data types with inbuilt conflict resolution.
- Transforms non-commutative operations to commutative operations with certain tricks
- (Assumption)
  - All operations will reach all nodes eventually.
- (Guarantee)
  - All nodes will converge to same state.

# Replicated Set

- Remove() does not commute with add()

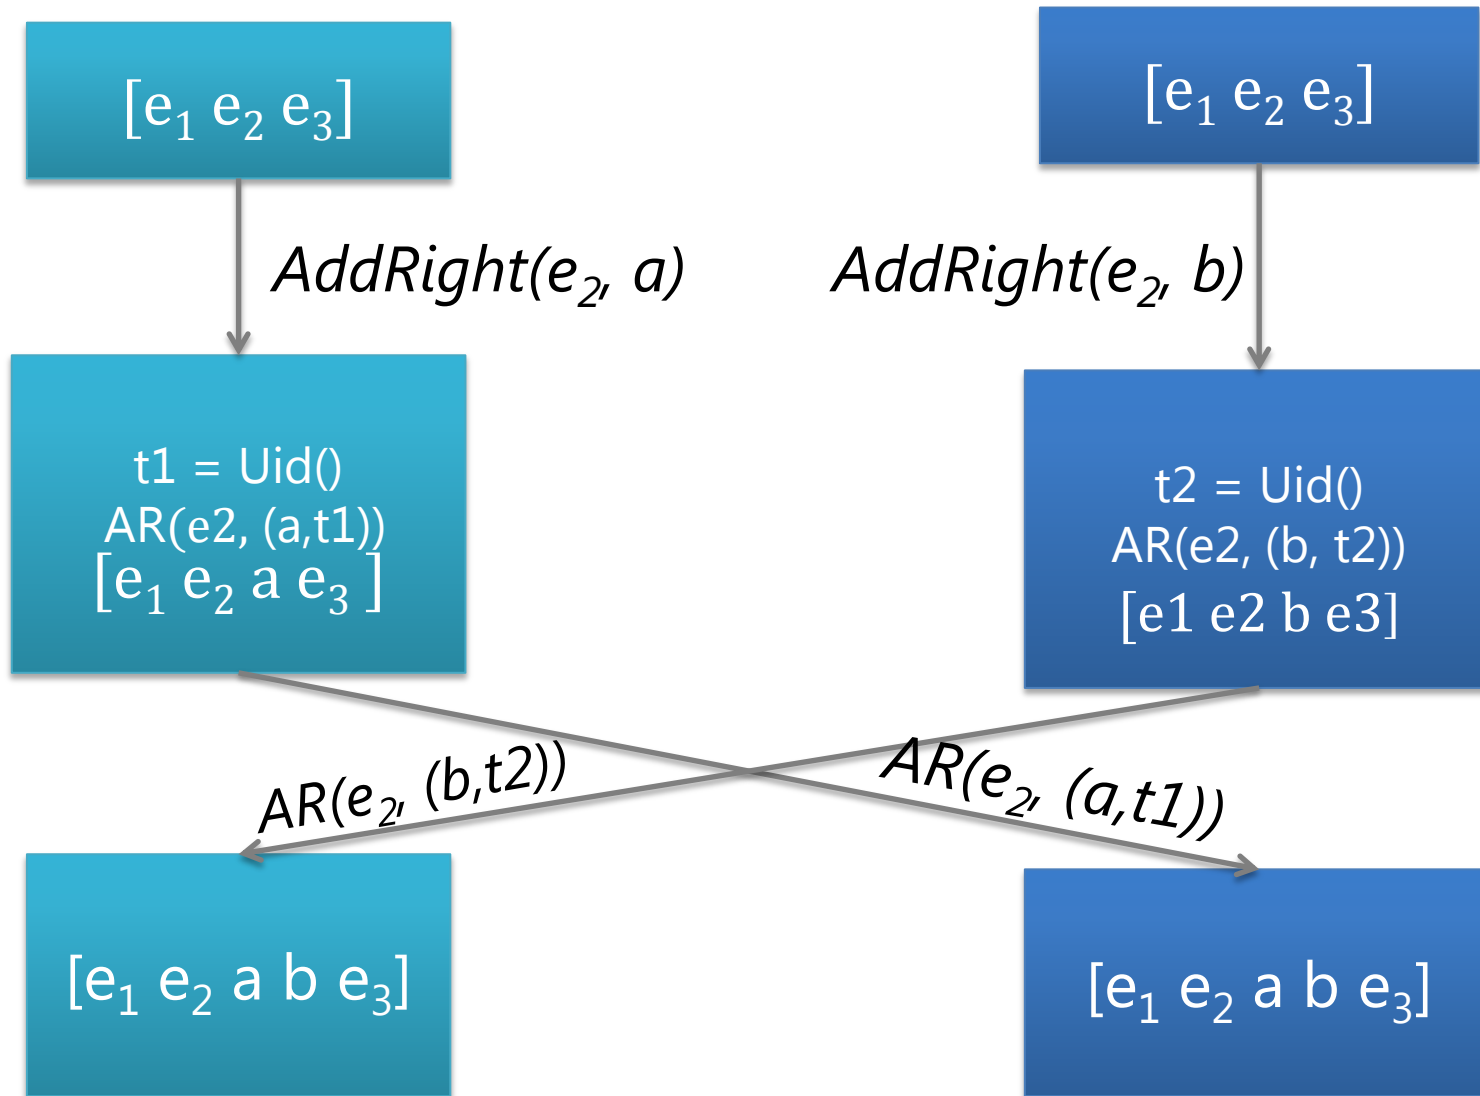


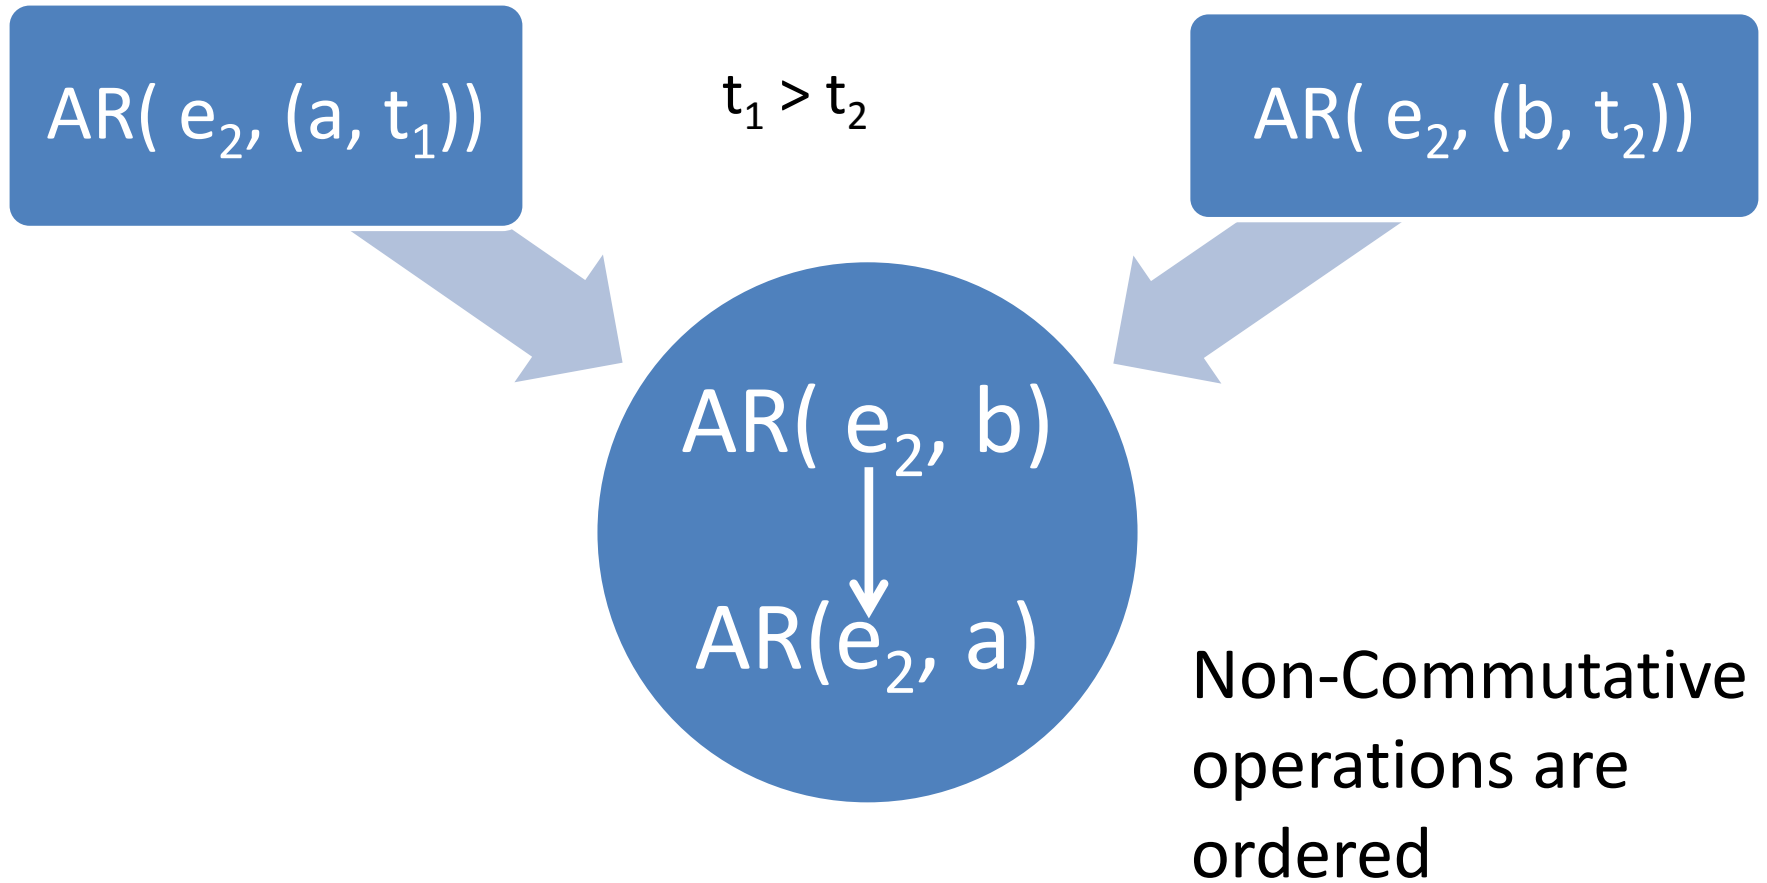
# Observed Remove set



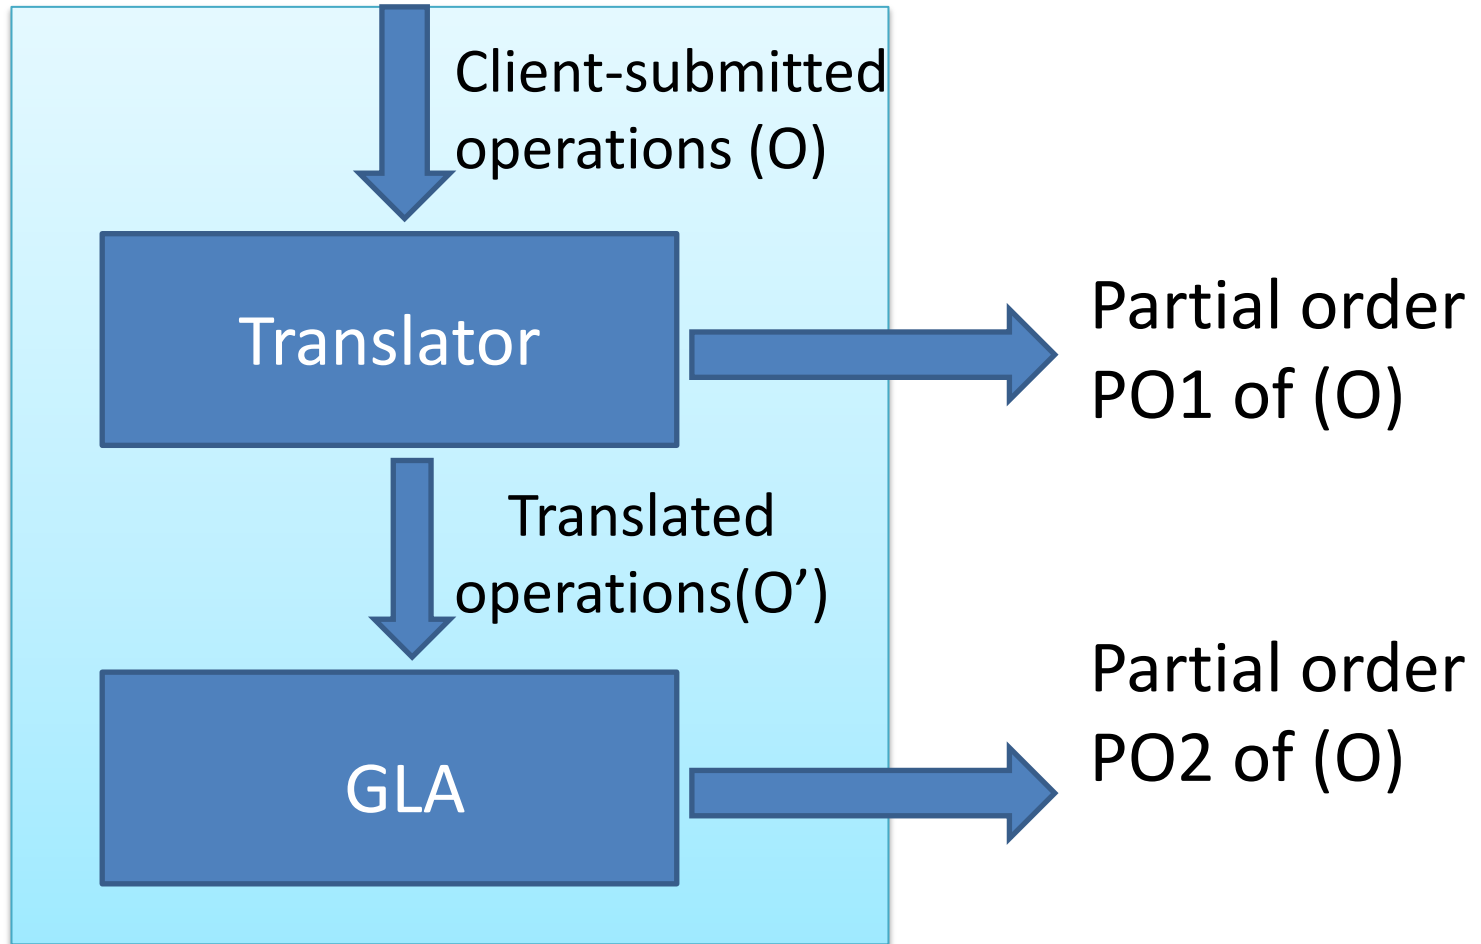
Orders non-commutative  
operations

# Replicated Growable array (RGA)





# Big-picture




What if PO1 conflicts PO2?


# Previous Example...

Sequence:  $[e_1 \ e_2 \ e_3]$

## GLA Ordering




## Witness 1



X


## Witness 2



X



# Problem!



- $op_1$  and  $op_2$  don't commute
- Deterministic ordering orders  $op_2$  *\_before\_*  $op_1$

- Effect:
  - $v_1$  exposes the effect of  $Op_1$ .
  - $Op_2 < Op_1$ ,
  - $v_1$  must also reflect the effect of  $op_2$  (to guarantee consistency across values  $v_1$  and  $v_2$ ).
- In general, this could be difficult!

What extra conditions  
we need for  
strong consistency?

# Solution


- Nullifying operations:
  - $Op1(Op2(s)) = Op1(s)$  where  $s$  is the state.
  - $Op1$  is said to nullify  $Op2$ .
- In set  $Add(x)$ ,  $Remove(x)$  have nullifying property.
  - $Add(x, Remove(x, s)) = Add(x, s)$
  - $Remove(x, Add(x, s)) = Remove(x, s)$
- Sequence doesn't have this.
  - $\{AddRight(e, x), AddRight(e, y)\} \neq \{AddRight(e, y)\}$
- **Sufficient condition:** If all pairs of non-commutative operations have nullifying property then strong consistency is guaranteed.

# Why nullifying works


Consider op1, op2 : two non-commutative operations.



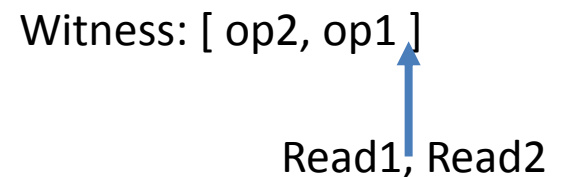

*Pre-deterministic  
ordering*



*GLA ordering1*



*GLA ordering2*



# Summary

- For data-types to have strong consistency
  - Commutative operations
    - Ordering decided by GLA.
  - Non-commutative operations
    - Some pre-deterministic ordering is decided.
    - All pairs of non-commutative operations should have nullifying semantics to obey this pre-defined order.

# Characterization

- Data-types with nullifying semantics
  - Set (2P set, OR-set)
  - Graph(2P2P based, OR-set based).
  - Directed Acyclic graph
  - Key-value pair
  - Leaderboard
- Data-types which fail
  - Sequences (AddRight())
- **Can we do better?**

# Thanks!

- References:
  - Generalized lattice agreement by *Jose, Sriram, Kaushik, Rama, Kapil*
  - CRDT by *Marc Shapiro*
  - Kapil's slides for MSR-summer school talk.
  - Windows 8 metro UI(Cover slide design)